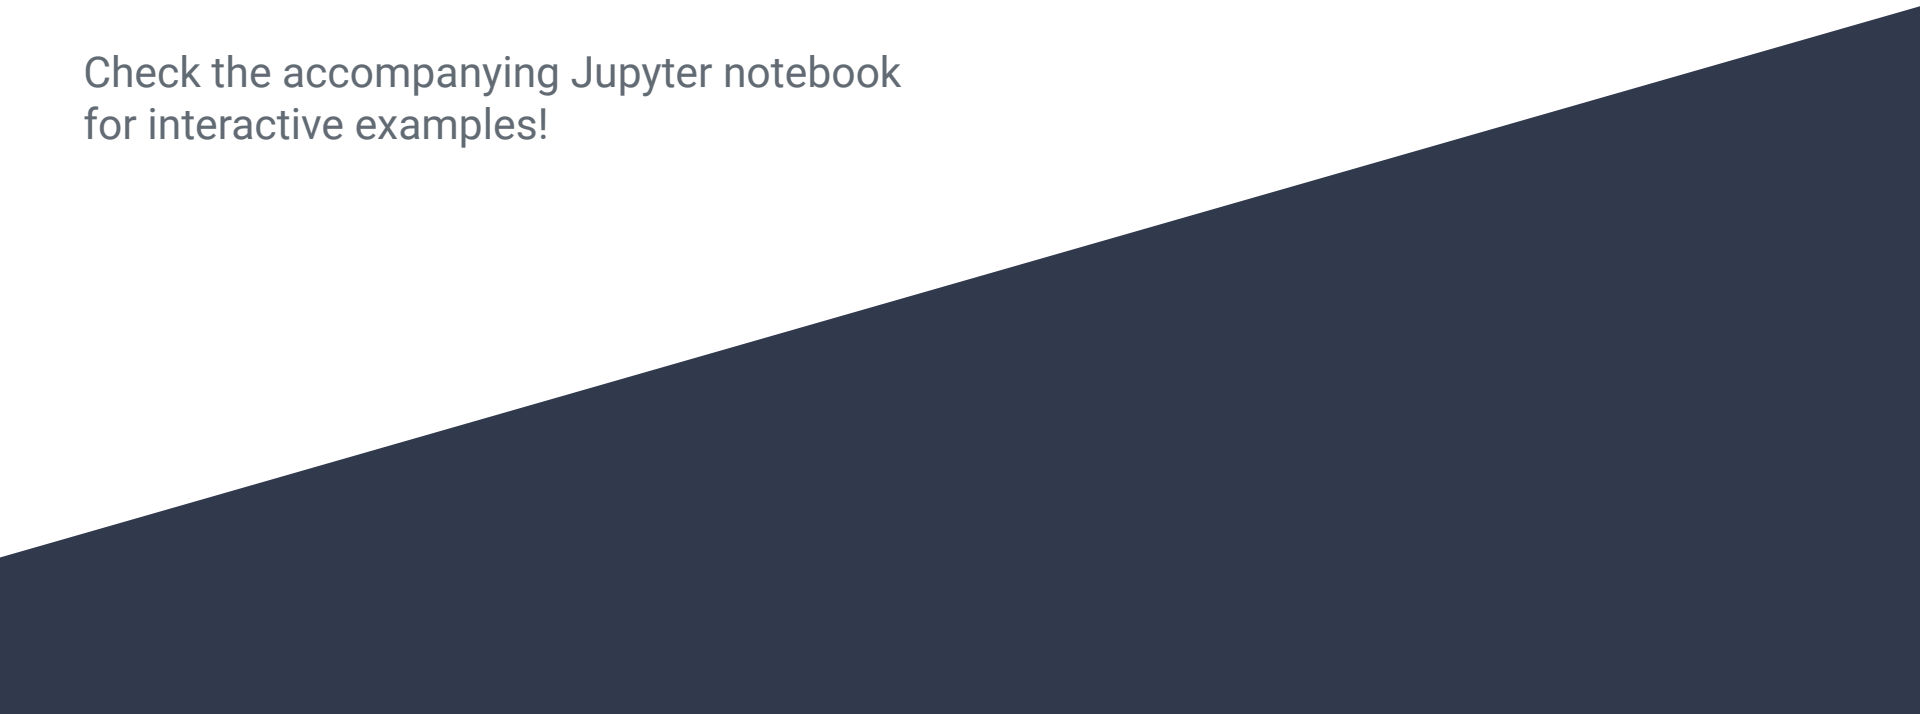


Week 5 - Advanced NumPy

Check the accompanying Jupyter notebook
for interactive examples!

A dark blue diagonal gradient bar that starts from the bottom left and extends towards the top right, covering the lower half of the slide.

Advanced Indexing

Advanced indexing (also sometimes called **fancy indexing**) happens when the selection object...

- ... is a non-tuple sequence object (e.g. a list)
- ... is a `np.ndarray`
- ... is a tuple, where at least one element is one of the above

Advanced indexing always returns a **copy** instead of a view.

```
array = np.arange(9).reshape(3, 3)

array[1, 2]           # basic indexing
array[(1, 2)]         # basic indexing

array[[1, 2]]         # fancy indexing

index = np.arange(3)
array[index]          # fancy indexing

array[[1, 2], 0]      # fancy indexing
```

Integer Array Indexing

Simple case:

Index is a tuple of **np.ndarrays** of shape **(n,)** with integer dtype, **one for each dimension**:

- selects **n** elements
- *i*-th entry of each array corresponds to index of the *i*-th element along that dimension

```
array = np.arange(9).reshape(3, 3)

index_row = np.array([0, 1, 2])
index_col = np.array([2, 2, 0])

array[index_row, index_col]
# → [2, 5, 6]

index = np.zeros(10, dtype=int)

array[index, index]
# → [0, 0, 0, 0, 0, 0, 0, 0, 0, 0]
```

Integer Array Indexing

Generalization:

Index is a tuple of np.ndarrays of **any broadcast-compatible shapes**, one for each dimension:

- selects as many elements as the broadcasted shape demands
- output shape is the broadcasted shape
- otherwise identical to the case before

```
array = np.arange(9).reshape(3, 3)

index_row = np.array([[0, 1], [0, 1]])
index_col = np.array([[0, 0], [1, 1]])

array[index_row, index_col]
# → [[0, 3],
#     [1, 4]]

array[index_row, 0]
# → [[0, 3],
#     [0, 3]]
```

Combining Advanced Indexing and Slicing

... is possible but can be complicated.

If you really need it, and you don't get the result you'd intuitively expect, read this:

<https://numpy.org/doc/stable/reference/arrays.indexing.html#combining-advanced-and-basic-indexing>

Special Case: Single Nested List

What do you think the example on the right will return?

```
array = np.arange(9).reshape(3, 3)  
array[[[0, 1], [0, 1]]]
```

Special Case: Single Nested List

What do you think the example on the right will return?

Answer: A FutureWarning!

Currently, this is interpreted as

```
array[[0, 1], [0, 1]]
```

i.e. two separate index lists. In a future version of NumPy, this will change to

```
array[np.array([[0, 1], [0, 1]])]
```

i.e. a single 2D array indexing only the first dimension.

⇒ **avoid using this for now.**

```
array = np.arange(9).reshape(3, 3)
array[[[0, 1], [0, 1]]]
```

current behaviour

→ [0, 4]

future behaviour

```
# → [[[0, 1, 2],
#      [3, 4, 5]],
#
#      [[0, 1, 2],
#      [3, 4, 5]]]
```

Example: `np.argsort`

How to sort one array according to values of another array?

```
letters = np.array(["y", "P", "n", "o", "t", "h"])
order = np.array([2, 1, 6, 5, 3, 4])

letters[np.argsort(order)]
# → ['P', 'y', 't', 'h', 'o', 'n']
```

We sort the array of numbers using `np.argsort`, which returns not the elements but their **indices**.

Because of **advanced indexing**, this array of indices can be used to index the `letters` array in the correct order.

np.nonzero & Friends

`array.nonzero()` returns the indices of all values in an array that are True-ish ($\neq 0$).

The output format is a tuple of **n** 1D arrays, where **n** is the number of dimensions of the array.

⇒ appropriate format for indexing

`np.where` and `np.argwhere` are similar, but have slightly different use-cases.

```
array = np.array([[0, 2], [0, -1]])
```

```
array.nonzero()
```

```
# → ([0, 1], [1, 1])
```

```
# often used with boolean arrays
```

```
(array < 0).nonzero()
```

```
# → ([1], [1])
```

```
# can be used for indexing
```

```
# (please immediately forget this)
```

```
array[(array < 0).nonzero()]
```

```
# → [-1]
```

Boolean Array Indexing

Boolean arrays can be used directly for indexing, without the need for `array.nonzero()`

If the array and the index have the same shape, a 1D array with all elements where the index is `True` is returned.

You can also use boolean indexing on individual dimensions. In conjunction with slices, this can again get a bit complicated.

```
array = np.array([[0, 2], [0, -1]])  
  
index = np.array([[True, False],  
                  [False, True]])  
  
array[index]      # → [0, -1]  
  
array[array < 0]  # → [-1]  
  
array[[True, False], 0]  
# → [0]
```

Sidenote: Bitwise Operators

If you want to perform element-wise logical operations (like and, or, xor) on NumPy arrays, you can use the bitwise operators:

and	&
or	
exclusive or	^
not	~

```
array = np.array([[0, 2], [0, -1]])  
  
array[(array < 0) | (array > 0)]  
# → [2, -1]
```

Concatenating Arrays

NumPy provides confusingly many functions for concatenating arrays, but they all build on `np.concatenate`.

It takes a tuple of arrays (of appropriate shape) and concatenates them along a given axis.

`np.r_` is a useful shorthand.

- does not use function parentheses but `[]`
- first element is a **string** giving the axis
- remaining elements are arrays to be concatenated

```
array = np.arange(4).reshape(2, 2)

np.concatenate((array, array), axis=0)
# → [[0, 1],
#     [2, 3],
#     [0, 1],
#     [2, 3]]

np.concatenate((array, array), axis=1)
# → [[0, 1, 0, 1],
#     [2, 3, 2, 3]]

np.r_["1", array, array]
# → [[0, 1, 0, 1],
#     [2, 3, 2, 3]]
```

Coordinate Grids

Grids allow you to get all possible pairs of two arrays. This is useful for example when evaluating a multidimensional function over a coordinate grid.

NumPy provides the `meshgrid` function:

- takes any number of 1D arrays (i.e. number of dimensions of the grid)
- returns a list of n n -dimensional arrays that constitute all possible combinations

There are also the `np.mgrid` and `np.ogrid` objects that do similar things.

```
x_values = np.array([2, 3, 4])
y_values = np.array([-1, 0])

xx, yy = np.meshgrid(x_values, y_values)

xx
# → [[2, 3, 4],
#     [2, 3, 4]]

yy
# → [[-1, -1, -1],
#     [ 0,  0,  0]]

some_2d_function(xx, yy)
```